ETH Zürich   Prof. Marc Pollefeys
Computer Science Department   Prof. Vittorio Ferrari
                                                  Prof. Luc Van Gool
                                                              HS 10

Computer Vision - Exercise 8

Shape Context

Shape Context Descriptors and Shape Matching

How can we enable a computer to recognize shapes the way that a human observer does? Looking at two handwritten letters, they appear very similar as images but will be very different if one compares the pixel intensity values. This exercise will introduce a feature descriptor called *shape context* [1] and explore their use in matching two shapes. You will then explore the use of shape contexts for shape classification.

8.1 Shape Matching (60%)

Your first task is to implement a shape matching algorithm using MATLAB. We will make use of a descriptor called the *shape context descriptor*, described in detail in [1], which is available in the released files as *BelongiePAMI02.pdf*.

Suppose we are given a set of points from a template contour for which we want to match to a set of points on a target contour. One possible algorithm to achieve this is:

a) Compute shape context descriptors for the points from both sets, the template and the target contour.

b) Estimate the cost matrix between the two sets of descriptors.

c) Use the cost matrix to solve the correspondence problem between the two sets of descriptors, finding the one-to-one matching that minimizes the total cost (e.g. with the provided Hungarian algorithm).

d) Use the solution of the correspondence problem to estimate a transformation from template to target points (e.g. with Thin Plate Splines) and perform this transformation on the template points.

e) Iterate steps (a-d).

We provide parts of the algorithm in `shape_matching.m`. Your task is to provide the missing code marked in `shape_matching.m` by comments. The MATLAB functions to be written are described below.

a) **Shape Context Descriptors (25%)**

Write a function which computes the shape context descriptors for a set of points. Your function should have the following form:

```
d = sc_compute(X,nbBins_theta,nbBins_r,smallest_r,biggest_r)
```

with the output `d` containing the shape context descriptors for all input points, and the inputs given by:

- set of points, `X`
- number of bins in the angular dimension, `nbBins_theta`
- number of bins in the radial dimension, `nbBins_r`
- the length of the smallest radius, `smallest_r`
- the length of the biggest radius, `biggest_r`

Hint: The shape context descriptor is described in detail in [1] in pages 511-513. For increased robustness, implement the normalization of all radial distances by the mean distance of the distances between all point pairs in the shape.

b) **Cost Matrix (10%)**

Write a function which computes a cost matrix between two sets of shape context descriptors. The cost matrix should be an $n \times m$ matrix giving the cost of matching two sets of points based on their shape context descriptors. One possibility is to use the χ^2 test statistic:

$$C_{gh} = \frac{1}{2} \sum_{k=1}^{K} \frac{[g(k) - h(k)]^2}{g(k) + h(k)} \tag{1}$$

where C_{gh} is the shape matching costs between two points with shape context descriptors g and h, each made up of $K = nbBins\_theta \times nbBins\_r$ bins.

The function should have the following form:

```
C = chi2_cost(s1,s2)
```

where the output `C` is the cost matrix for matching two sets of shape context descriptors `s1` and `s2`.

c) **Hungarian Algorithm**

We provide the code for the Hungarian algorithm which performs a one-to-one matching of the points based on the cost matrix, minimizing the total cost. This code is provided in the released file:

`hungarian.m`

d) Thin Plate Splines (25%)

From the point correspondences, we can estimate a plane transformation $T : \mathbb{R}^2 \to \mathbb{R}^2$ that maps any point (not only those originally sampled) from one shape to the other. We will use a thin plate spline (TPS) model:

$$f(x,y) = a_1 + a_x x + a_y y + \sum_{i=1}^{n} \omega_i U(\| (x_i, y_i) - (x,y) \|) \tag{2}$$

where $U(t) = t^2 log(t^2)$ and $U(0) = 0$ and it holds that $\sum_{i=1}^{n} \omega_i = \sum_{i=1}^{n} \omega_i x_i = \sum_{i=1}^{n} \omega_i y_i = 0$.

This model gives a 1D output. Since we are interested in a 2D warping, we will use **two independent TPS models**, which we call f_x and f_y, to model the x and y coordinate transformations respectively. These combine to give the full transformation T:

$$T(x,y) = (f_x(x,y), f_y(x,y)) \tag{3}$$

For both of these models, (x_i, y_i) are given by the points coordinates in the **original** template shape, and the appropriate coordinates of the corresponding points on the target shape give the values v_i used to solve for the TPS model.

For each TPS model, you will have to solve a system of the form

$$\begin{pmatrix} K + \lambda I & P \\ P^T & 0 \end{pmatrix} \begin{pmatrix} \omega \\ a \end{pmatrix} = \begin{pmatrix} v \\ 0 \end{pmatrix} \tag{4}$$

where $K_{ij} = U(\| (x_i, y_i) - (x_j, y_j) \|)$, the i-th row of P is $(1, x_i, y_i)$, ω and v are column vectors formed from ω_i and v_i, respectively, a is the column vector with elements a_1, a_x, a_y, and λ is a regularizer (see [1] for details).

Solving this system allows not only the transformation to be determined, but the bending energy of the transformation to be computed. This bending energy can be used as a measure of the cost of the shape matching. It is given by:

$$E = \omega^T K \omega \tag{5}$$

Your task is to implement a MATLAB function that computes the weights ω_i and a_1, a_x, a_y for both f_x and f_y. The function should have the following form:

```
[w_x w_y E] = tps_model(X,Y,lambda)
```

where the outputs `w_x` and `w_y` are the parameters (ω_i and a_i) in the two TPS models, E is the total bending energy and the inputs are as following:

- points in the template shape, `X`
- corresponding points in the target shape, `Y`
- regularization parameter, `lambda`

You can use a MATLAB solver for the linear system (if you have to find x, in $Ax = b$, you can find the solution in MATLAB with $x = A\backslash b$). Having w_x and w_y you are now able to perform the transformation (warping) on the template points.

Hint: For regularization, set *lambda* to the square of the mean distance between two target points.

8.2 Shape Classification (40%)

Your second task is to use your code from the first task to perform image classification based on shape using a *k*-nearest-neighbour classifier. The images to be used can be found in the released file:

```
dataset.mat
```

Load the objects with the `load` command in MATLAB. This file contains a structure array called `objects`. Each structure in the array contains the image, the extracted edge points and the class name for a single object. There are 5 images for each category (*heart*, *fork* and *watch*).

The following steps must be carried out for the *k*-nearest neighbour classification:

a) Determine shape matching costs between a test shape and all training shapes.

b) Classify the test shape based on the labels of the *k*-nearest neighbour training shapes.

You will evaluate this classifier with leave-one-out cross: for a dataset of size n, $n - 1$ samples are used for training, while the n^{th} sample is used for testing. This procedure is repeated n times so that all samples are used for testing once.

For this task, we provide parts of the algorithm in `shape_classification.m`. Your task is to provide the missing code marked in `shape_classification.m` by comments. The MATLAB functions to be written are described below.

a) **Shape Matching (25%)**

For each shape, a set of points lying on the edges is provided. Write a function that randomly samples n points from this list. The function should have the following form:

```
X_nsamp = get_samples(X,nsamp)
```

where the output X_nsamp is of size *nsamp* $\times$ 2 and contains the cooordinates of the sampled points, given the points X and the number of samples to take `nsamp`. Use this function and the shape matching algorithm already implemented (function `shape_matching.m`) to compute the cost of matching each test shape to all training shapes. Sampling 100 points for each shape should be sufficient.

The function that computes the cost of matching each test shape to all training shapes should have the following form:

```
matchingCostMatrix = compute_matching_costs(objects,nsamp)
```

where the output is a square matrix (15×15 for our problem) containing the shape matching costs and the input of this function is given by the structure *objects* and the number *nsamp* of points that are sample from each shape (100 for our problem). The (i,j) element of the *matchingCostMatrix* represents the shape matching cost of matching the test shape *i* to the training shape *j*. The diagonal of this matrix should have elements of infinite energy. Building this matrix with this form allows for easy *k*-nearest neighbour classification with leave-one-out cross validation.

b) Nearest-Neighbour Classifier (15%)

For a test shape, find the *k*-nearest neighbour matches in the training shapes, and use these to assign a label to the test shape.

You should write a MATLAB function with the following form:

```
testClass = nn_classify(matchingCostVector,trainClasses,k)
```

Your function should take as its input the shape matching costs obtained by matching the test shape to all training shapes `matchingCostVector`, the class labels of the training shapes `trainClasses`, and the number of neighbours to consider *k*. It should return the class label of the test shape as its output `testClass`.

Warning: It takes several minutes (~ 30) to run the classification for all the objects. Be sure to allocate enough time!

8.3 Hand In

Write up a short report explaining the main steps of your implementation and discussing the results of the methods. Make sure to include answers to the following questions:

- Is the shape context descriptor scale-invariant? Explain why or why not.

- What is the average accuracy of your classifier?

- How does your classification accuracy vary with the number of neighbours *k*?

- If instead of your own sampling function, you use the one that we provide (`get_samples_1.m`), do you get better classification results? Why or why not?

Send the report together with the source code for your implementation (including the following functions `shape_matching`, `sc_compute`, `chi2_cost`, `tps_model`, `get_samples`, `compute_matching_costs` and `nn_classify`) to the Moodle exercise page. If you have any questions about this exercise, please email to `yawei.li@vision.ee.ethz.ch`.

References

[1] Belongie, S., Malik, J. and Puzicha, J., 2002, 'Shape Matching and Object Recognition Using Shape Contexts', IEEE Trans. PAMI